

Pipeline and Goals

1. Each student implements a **generative model for tabular data**.
2. Their model should:
 - take **raw data**,
 - **discretize** (or otherwise preprocess) it,
 - **generate synthetic tabular data**,
 - **evaluate quality via TSTR** (train on synthetic, test on real),
 - **save outputs + metrics to disk**.
3. All students must share:
 - the **same train/test split**,
 - the **same preprocessing tools** (for fairness),
 - the **same evaluation pipeline and metrics**.

You will implement **your own generative model for tabular data**.

Your model will:

1. Take a **raw tabular dataset** (CSV).
2. **Discretize / preprocess** it using the **shared tools**.
3. **Train a generative model** on the training split.
4. **Generate synthetic tabular data**.
5. Evaluate the synthetic data by:
 - training classifiers on synthetic data,
 - testing on **real** data (TSTR: Train on Synthetic, Test on Real),
 - saving all outputs & metrics to disk.

All students **must use the same pipeline and utilities** so results are comparable.

Codebase Components You **MUST** Use

Below is the pipeline you must follow, with exact functions + file locations.

1- Preprocessing / Discretization (from RAW → DISCRETE CSV)

You are NOT allowed to write your own discretizer.

You must use the shared function.

- **What to use**

- `discretize_preprocess(file_path, output_path, bins=10, strategy='uniform')`
- **File:** `utils.py`

- **What it does**

- Loads raw CSV (`file_path`).
- Treats '?' as missing and drops all-NaN columns.
- Splits into:
 - $X = \text{all columns except the last,}$
 - $y = \text{last column (assumed to be the target/label).}$
- For **numerical features**:
 - converts to numeric,
 - fills missing values with the column median,
 - discretizes into integer bins using KBinsDiscretizer.
- For **categorical features**:
 - fills missing with 'Missing',
 - encodes with LabelEncoder.
- For the **target column**:
 - fills missing with 'Missing',
 - encodes with LabelEncoder.
- Saves a fully **discrete integer** dataset to `output_path`.

So your first step is always something like:

```
from utils import discretize_preprocess

discretize_preprocess(
    file_path="raw_data/mydataset.csv",
    output_path="discretized_data/mydataset.csv",
    bins=10,
    strategy="uniform",
)
```

Preprocessing / discretization

- Discretize raw data
 - `discretize_preprocess(file_path, output_path, bins, strategy)`
→ file: `utils.py`
 - Takes raw CSV, cleans it, discretizes numeric columns into integer bins, label-encodes categoricals & target, and saves a **fully discrete** dataset.
- Alternate continuous+encoded preprocessing
 - `encode_preprocess(file_path, output_path)`
→ file: `utils.py`
 - Cleans numeric columns (no binning), encodes categoricals, encodes target.

2- Train/Test Split (SHARED FOR EVERYONE)

All students must use the **same split** logic.

- What to use
 - `split_dataset(input_csv, output_dir, test_size=0.2, seed=42, ...)`
 - File: `split_dataset.py`
- How it's invoked in the pipeline

The function is called inside:

- `TrainTestSplitPipeline.run(...)`
 - File: `katabatic.pipeline.train_test_split.pipeline`
(you saw this earlier as `TrainTestSplitPipeline`)
- What it does
 - Loads the **preprocessed** CSV (`input_csv`).
 - Runs a **stratified** train/test split on the **last column** (the label).
 - Saves:
 - `train_full.csv`, `test_full.csv`
 - `x_train.csv`, `y_train.csv`
 - `x_test.csv`, `y_test.csv`
 - All in `output_dir`.

You must not change this function or splitting logic.

3- Your Generative Model (STUDENT CODE)

Generative model example (CoDi)

- `class CODI(Model)`
 - file: `codi.py`
 - `train(dataset_dir, synthetic_dir, **kwargs)`:
 - loads `x_train.csv` (+ optional `y_train.csv`) from `dataset_dir`,
 - infers schema, encodes data (using CoDi-specific helpers),
 - builds diffusion models,
 - trains them,
 - generates synthetic samples,
 - saves `x_synth.csv` and `y_synth.csv` into `synthetic_dir`,
 - saves `schema.json` and `metadata.json`.
 - Low-level utilities for CoDi
 - file: `codiutils.py`
 - `infer_schema`, `encode_dataframe`, `decode_dataframe`, `TabularUNet`, `GaussianDiffusionTrainer`, `GaussianDiffusionSampler`, `MultinomialDiffusion`, etc.

- Where to look for an example
 - `class CODI(Model)`
 - File: `codi.py`
- What your `train(...)` MUST do

Given:

- `dataset_dir` containing:
 - `x_train.csv`
 - `y_train.csv` (optional but recommended)
- `synthetic_dir` = directory to write your synthetic data

Your `train(...)` must:

1. Load the training data:

```
python

x_train_path = os.path.join(dataset_dir, "x_train.csv")
y_train_path = os.path.join(dataset_dir, "y_train.csv") # if used
```

2. Fit your generative model on this training data.
3. Generate **synthetic samples** with the **same feature structure**.
4. Save to **these exact file names**:

- `synthetic_dir/x_synth.csv`
- `synthetic_dir/y_synth.csv`

(column order and label name must be consistent with the original.)

- Optional helpers you can reuse (from CoDi)

From `codiutils.py`:

- `infer_schema(df, categorical_threshold=20)`
- `encode_dataframe(df, schema)`
- `decode_dataframe(df_encoded, schema)`
- `TabularUNet`, `GaussianDiffusionTrainer`, `GaussianDiffusionSampler`,
`MultinomialDiffusion`, `get_device`, `set_global_seed`

You can completely ignore CoDi internals and implement a different generative model (e.g., CTG, Gaussian copula, simple noise model) **as long as** you:

- obey the `Model` interface,
- read from `x_train.csv` / `y_train.csv`,
- write `x_synth.csv` / `y_synth.csv` in the same style.

4- Shared Pipeline to Run Everything

TSTR evaluation

- `class TSTREvaluation(Evaluation)`
→ file: `evaluation.py`
 - `load_data(synthetic_dir, real_test_dir)` loads:
 - synthetic: `x_synth.csv`, `y_synth.csv`
 - real: `x_test.csv`, `y_test.csv`
 - `evaluate()`:
 - Trains **four classifiers** on synthetic data:
 - LogisticRegression
 - MLPClassifier
 - RandomForestClassifier
 - XGBClassifier
 - Evaluates them on **real test data**.
 - Metrics:
 - Accuracy
 - F1 Score (weighted)
 - AUC (if binary labels)
 - Saves results to:
 - `Results/{dataset_name}/{model_name}_tstr.csv`

All the infrastructure needed for students to plug in their own generative models and still share:

- the same preprocessing (if they all call the same `utils.py` functions),
- the same train/test split (`split_dataset.py`),
- the same evaluation pipeline & metrics (`TSTREvaluation`),
- the same overall pipeline orchestration (`TrainTestSplitPipeline`).

constraints / interface

Students must implement a class like:

- `class MyModel(Model):`
 - Must have:
`def train(self, dataset_dir: str, synthetic_dir: str, **kwargs) -> "MyModel":`
 - Inside `train`, they must:
 - read `x_train.csv` (+ `y_train.csv` if needed) from `dataset_dir`,
 - train their generative model,
 - generate synthetic data,
 - write to:
 - `synthetic_dir/x_synth.csv`
 - `synthetic_dir/y_synth.csv`
- (same format as CoDi)

If they follow **exactly this interface**, they can be dropped into:

```
python

pipeline = TrainTestSplitPipeline(model=lambda: MyModel(...))
pipeline.run(
    input_csv=preprocessed_csv,
    output_dir=output_dir,
    synthetic_dir=synthetic_dir,
    real_test_dir=output_dir,
)
```

2. Where discretization happens

Right now, discretization is done **outside** the pipeline (e.g., in the notebook):

```
python

from utils import discretize_preprocess

dataset_path = ROOT / "raw_data" / "car.csv"
output_path = ROOT / "discretized_data" / "car.csv"
discretize_preprocess(str(dataset_path), str(output_path))
```

Note: Before running the pipeline, preprocess raw data using `discretize_preprocess` from `utils.py`

Extended pipeline overview (with pointers to code)

Here's a checklist, showing each required step and exactly which function/file to use.

Step 1 — Discretize raw data

Requirement: “Discretizes data”

- Use:
`discretize_preprocess(file_path, output_path, bins=10, strategy='uniform')`
- Found in:
`utils.py`
- What it does:
 - Loads raw CSV.
 - Cleans missing values.
 - Discretizes numeric columns into integer bins using KBinsDiscretizer.
 - Label-encodes categoricals and target.
 - Saves a fully discrete CSV where:
 - All feature columns are integers.
 - The last column is the label.

Students should not write their own discretizer — they should call this one.

Step 2 — Train/test split (shared for all models)

Requirement: “same train/test split”

- Use (indirect):
`split_dataset(input_csv, output_dir, test_size=0.2, seed=42, ...)`
- Found in:
`split_dataset.py`
- Usually invoked by:
`TrainTestSplitPipeline.run(...)`
→ file: (your earlier TrainTestSplitPipeline file)
- What it does:
 - Loads the preprocessed (discrete) CSV.
 - Performs a single stratified split (default 80/20).

- Saves:
 - train_full.csv, test_full.csv
 - x_train.csv, y_train.csv
 - x_test.csv, y_test.csv

Every model sees the same train/test indices if they all use the same input_csv and seed.

Step 3 — Student generative model

Requirement: “Each student implements a generative synthetic data model”

- Base class:
Model
→ file: katabatic.models.base_model (imported in codi.py)
- Example implementation to follow:
class CODI(Model)
→ file: codi.py
- Required behavior for student models:
 - Implement train(self, dataset_dir: str, synthetic_dir: str, **kwargs) that:
 1. Reads training data from:
 - dataset_dir/x_train.csv
 - Optionally dataset_dir/y_train.csv
 2. Trains their generative model.
 3. Generates synthetic data with the same column structure.
 4. Saves:
 - synthetic_dir/x_synth.csv
 - synthetic_dir/y_synth.csv
(label column name should match the original label name)

If students want to use a CoDi-like encoding/decoding for mixed-type data, they can reuse:

- infer_schema, encode_dataframe, decode_dataframe
→ file: codiutils.py

But that’s optional; the assignment might allow any model as long as it writes x_synth/y_synth in the same format.

Step 4 — Use the shared pipeline

Requirement: “all models should use the same pipeline”

- Use:
`TrainTestSplitPipeline(model=...)`
→ file: your TrainTestSplitPipeline file (shown earlier)
- Key behavior inside run(...):
 1. Calls `split_dataset(input_csv, output_dir, ...)`
→ file: split_dataset.py
 2. Instantiates model:
`current_model = self.model()`
 3. Trains the generative model and generates synthetic data:
`current_model.train(output_dir, synthetic_dir=synthetic_dir, ...)`
 - dataset_dir argument here is output_dir (where x_train/y_train live)
 - synthetic_dir is where the model must write x_synth.csv, y_synth.csv.
 4. Runs all evaluations in _evaluations, which includes TSTREvaluation.

Step 5 — TSTR evaluation (same for all models)

Requirement: “same evaluation pipeline / same metrics”

- Use:
`TSTREvaluation(synthetic_dir, real_test_dir)`
→ file: evaluation.py
(automatically used inside TrainTestSplitPipeline)
- It uses:
 - `x_synth.csv, y_synth.csv` from synthetic_dir
 - `x_test.csv, y_test.csv` from real_test_dir
- Models trained:
 - Logistic Regression (LR)
 - MLPClassifier (MLP)
 - RandomForest (RF)
 - XGBoost (XGBoost)
- Metrics computed:
 - Accuracy

- F1 Score (weighted)
- AUC (if binary classification)
- **Saved outputs:**
 - CSV file:
Results/{dataset_name}/{model_name}_tstr.csv
created by save_results_to_csv(results, synthetic_dir) in evaluation.py.

This ensures **every student** is judged by exactly the same TSTR protocol.

Here's the full pipeline map:

- **Discretizes data**
→ discretize_preprocess(file_path, output_path, bins, strategy)
file: utils.py
- **Splits into train/test (same for everyone)**
→ split_dataset(input_csv, output_dir, test_size=0.2, seed=42)
file: split_dataset.py
(called from TrainTestSplitPipeline.run)
- **Defines model interface all students must follow**
→ class Model (base)
file: katabatic.models.base_model
→ Example: class CODI(Model)
file: codi.py
- **Implements one generative model (CoDi example)**
→ class CODI(Model) with train(...), sample(...)
file: codi.py
- **Low-level building blocks for CoDi (optional reuse)**
→ infer_schema, encode_dataframe, decode_dataframe, TabularUNet,
GaussianDiffusionTrainer, GaussianDiffusionSampler, MultinomialDiffusion, get_device,
set_global_seed
file: codiutils.py
- **Common training + evaluation pipeline**
→ class TrainTestSplitPipeline(Pipeline)
file: your pipeline file (imports split_dataset, TSTREvaluation)

- **TSTR evaluation with fixed models & metrics**
→ class TSTREvaluation(Evaluation)
file: evaluation.py
- **Saving all outputs & metrics to disk**
 - Preprocessed data:
 - discretize_preprocess → saves preprocessed CSV
 - **file:** utils.py
 - Train/test splits:
 - split_dataset → saves train_full.csv, test_full.csv, x_*, y_*
 - **file:** split_dataset.py
 - Synthetic data:
 - CODI.train (and student models) → saves x_synth.csv, y_synth.csv
 - **file:** codi.py or their own model file.
 - Evaluation metrics:
 - TSTREvaluation.save_results_to_csv → saves {model_name}_tstr.csv
 - **file:** evaluation.py